

Joltik:

Enabling Energy-Efficient “Future-Proof” Analytics on Low-Power Wide-Area Networks

Mingran Yang, Junbo Zhang, Akshay Gadre,
Zaoxing Liu, Swarun Kumar, and Vyas Sekar


Carnegie Mellon University

LPWAN provides opportunity for many applications

- ✓ Long range
- ✓ Low power



Low-power wide-area network (LPWAN) Technologies



Industrial monitoring




Smart city



Smart farming

Generality vs. Power trade-off for LPWAN applications



Downstream Analytics

Energy generated →

Power outages →


Weather events →

..... →

- Occasional short samples
- Pre-determined summary statistics

New tasks → additional energy overhead

System design goal 1: generality



Downstream Analytics

Energy generated

Power outages


Weather events

.....



Support estimation on multiple
statistics simultaneously

System design goal 2: high-fidelity



Downstream Analytics

Energy generated


Power outages

Weather events

.....

↓
Error rate < 5%

System design goal 3: energy-efficiency



Downstream Analytics

Energy generated

Power outages

Weather events

.....

Existing solutions and limitations

Approach	Energy-efficiency	High-fidelity	Generality
Sub-sampling	✓	✗	✓
Lossless compression	✗	✓	✓
Lossy compression	✓	✗	✓
Sparse recovery	✓	✓	✗
Data-centric aggregation	✓	✓	✗


Joltik vs. existing solutions

Approach	Energy-efficiency	High-fidelity	Generality
Sub-sampling	✓	✗	✓
Lossless compression	✗	✓	✓
Lossy compression	✓	✗	✓
Sparse recovery	✓	✓	✗
Data-centric aggregation	✓	✓	✗
Joltik	✓	✓	✓

↑
Universal sketching

- Support the estimation on broad class of statistics
- No requirements on raw data / network topology

Joltik vs. existing solutions



Outline for this talk

- Motivation
- System workflow and challenges
- System design
 - Low memory footprint - reducing memory footprint
 - Low power - reducing communication footprint
 - Low CPU - reducing computation overhead
- Implementation and evaluation
- Conclusions

Outline for this talk

- Motivation
- System workflow and challenges
- System design
 - Low memory footprint - reducing memory footprint
 - Low power - reducing communication footprint
 - Low CPU - reducing computation overhead
- Implementation and evaluation
- Conclusions

Joltik system workflow

Solar sensor



Pre deployment
Configuration



Base station


Sensing

Raw sensed data

Practical realization of
universal sketching

Sketch summaries

Wireless
Communication



Query:
..... (*new statistics*)

Sketch summaries

Background on universal sketching

- Sketching algorithm:

Input Data Stream



Sketches



Query to estimate
certain statistics

e.g., count-min sketch serves as a **frequency table** of events in a data stream

- Universal Sketching:

Input Data Stream



Universal
Sketches



Query to estimate
multiple statistics


Why universal sketching is a promising solution?

Universal sketching can estimate **many already known (generality)**, or even **possibly unforeseen (future-proof)** statistics at the same time.

Estimation tasks in solar sensor	Statistics supported by universal sketching
Energy generated	L1-norm
Power outages	Zero-draw time
Weather events	Change
Anomaly detection	Entropy
Voltage volatility	L2-norm
.....


Table: Statistics relevant to solar farm

Data structure in universal sketching algorithm




Universal Sketches


Data structure in universal sketching algorithm



How sensed value updates a single layer



How sensed value updates the entire structure



Joltik system challenges

- Challenge 1: Reduce memory footprint of universal sketches

Native Universal Sketching

← Hundreds of KBs memory

e.g., 300 KB to achieve 95% accuracy on a real-world dataset

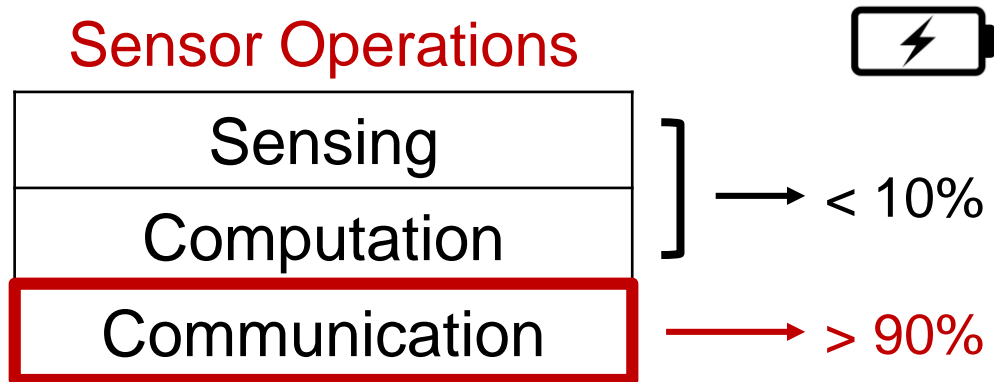


← Limited on-chip memory

e.g., the sensor board in our testbed only has 128 KB memory

Joltik system challenges

- Challenge 2: Reduce communication footprint



Joltik system challenges

- Challenge 3: Reduce computation overhead of sketch update

Universal sketch operations

Hash computations
Arithmetic calculations
Counter updates
Heap updates
.....


Outline for this talk

- Motivation
- System workflow and challenges
- System design
 - Low memory footprint - reducing memory footprint
 - Low power - reducing communication footprint
 - Low CPU - reducing computation overhead
- Implementation and evaluation
- Conclusions

Outline for this talk

- Motivation
- System workflow and challenges
- System design
 - Low memory footprint - reducing memory footprint
 - Low power - reducing communication footprint
 - Low CPU - reducing computation overhead
- Implementation and evaluation
- Conclusions


Reducing memory footprint: “Inverted Pyramid”



Gradually reduce
number of columns

Insight of “Inverted Pyramid” structure


- Lower layers need to handle much smaller number of samples




Outline for this talk

- Motivation
- System workflow and challenges
- System design
 - Low memory footprint - reducing memory footprint
 - Low power - reducing communication footprint
 - Low CPU - reducing computation overhead
- Implementation and evaluation
- Conclusions

Reducing communication footprint



Data packet

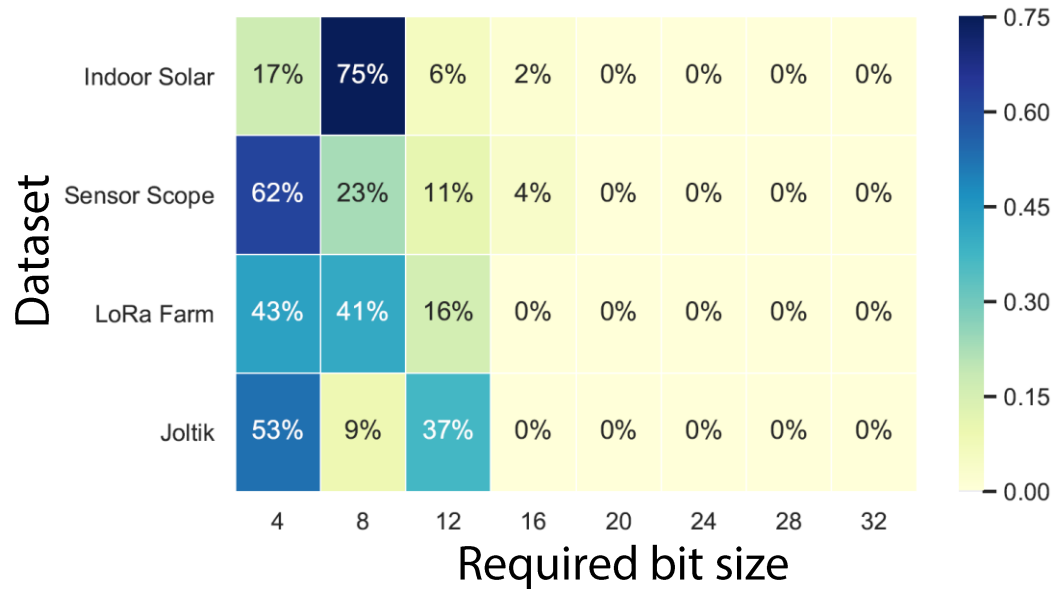


Observation on sketch counters

- Only small part of counters have large values

5	3	7930	9	11
3	4	30	7	10
9	184	11	8	4
7	9	5	2859	3


A sketch counter example



Lossless encoding of sketch counters

Required Bitsize	Prefix	Bitsize before/after compress
4	00	16 / 6
8	01	16 / 10
12	10	16 / 14
16	11	16 / 18

Example:



All counters are originally 16 bits before encoding

Efficiently transmitting heavy-hitters heap

Key	Value
910	35552
804	34409
983	29374
905	19395
827	18890
...	...

Transmit ←

→ Reconstruct at base station

A heavy-hitter heap example

Outline for this talk


- Motivation
- System workflow and challenges
- System design
 - Low memory footprint - reducing memory footprint
 - Low power - reducing communication footprint
 - Low CPU - reducing computation overhead
- Implementation and evaluation
- Conclusions

Bottleneck analysis



- Top two CPU performance bottlenecks:

(a) hash computations

(b) counter updates



Reduced sketch update



Intuition: less samples in lower layers → less collisions → more accurate result

End-to-end deployment

- User input:
 - (a) Data collection rate - Rcl
 - (b) Period of transmission - T
- Parameter tuned by Joltik:
 - (a) Sketch Size - S



Rcl	T	S	Lifetime	Accuracy
10 Hz	1 day	30 KB	2751 days	$94.30\% \pm 3.27\%$
10 Hz	1 day	60 KB	1613 days	$96.61\% \pm 2.25\%$
10 Hz	1 day	90 KB	1141 days	$97.50\% \pm 1.92\%$

Table: Joltik deployment example

Outline for this talk

- Motivation
- System workflow and challenges
- System design
 - Low memory footprint - reducing memory footprint
 - Low power - reducing communication footprint
 - Low CPU - reducing computation overhead
- Implementation and evaluation
- Conclusions


Joltik implementation and evaluation setup




Real world testbed in a campus building at CMU

- Microcontroller: NUCLEO-L476RG
- Sensor board: X-NUCLEO-IKS01A2
- RF Frontend: SX1276 LoRa Transceiver

Joltik provides better energy-accuracy trade-off for “future-proof” analytics



Joltik supports multi-task handling and provides generality



(Positive values imply Joltik is worse and vice versa)

AppSet1 = {Cardinality}
AppSet2 = {Cardinality, Entropy}

AppSet3 = {Cardinality, Entropy, L2}
AppSet4 = {Cardinality, Entropy, L2, HH}

Summary

- Goal

- Build a general, accurate, and energy-efficient sensor analytics framework

- Our approach

- Propose a novel architecture by leveraging universal sketching
- Low memory footprint, low power and low CPU realization of universal sketching

- Joltik: Enabling energy-efficient “future-proof” analytics on LPWAN

- Support general and “future-proof” analytics
- Guarantee >5-year sensor battery life and <5% error rate on a range of statistics

<https://github.com/Joltik-project/Joltik>

Mingran Yang mingrany@andrew.cmu.edu